

華中科技大學

大数据处理实验报告

实验三：MapReduce 的基本操作

姓 名：刘日星

学 号：X2020I1007

院 系：计算机科学与技术

专 业：计科与金融

年 级：CS1804-交换生

指导教师：石宣化

2021 年 12 月 15 日

一：实验目的

- 1、了解 MapReduce 的用途
- 2、掌握 MapReduce 的基本命令

二：实验要求

1、第四节中的实验内容要附上完整的**实验过程截图以及必要的文字说明**，每个人的 IP 地址等不同，不能直接套用样例的截图。

2、请同学们在完成报告后，将报告的 pdf 版本命名为：**大数据实验三+姓名+学号.pdf**，并在这周五前,发到邮箱:

aaaaltaaaa@126.com(石老师班)

798792873@qq.com(郑老师班)

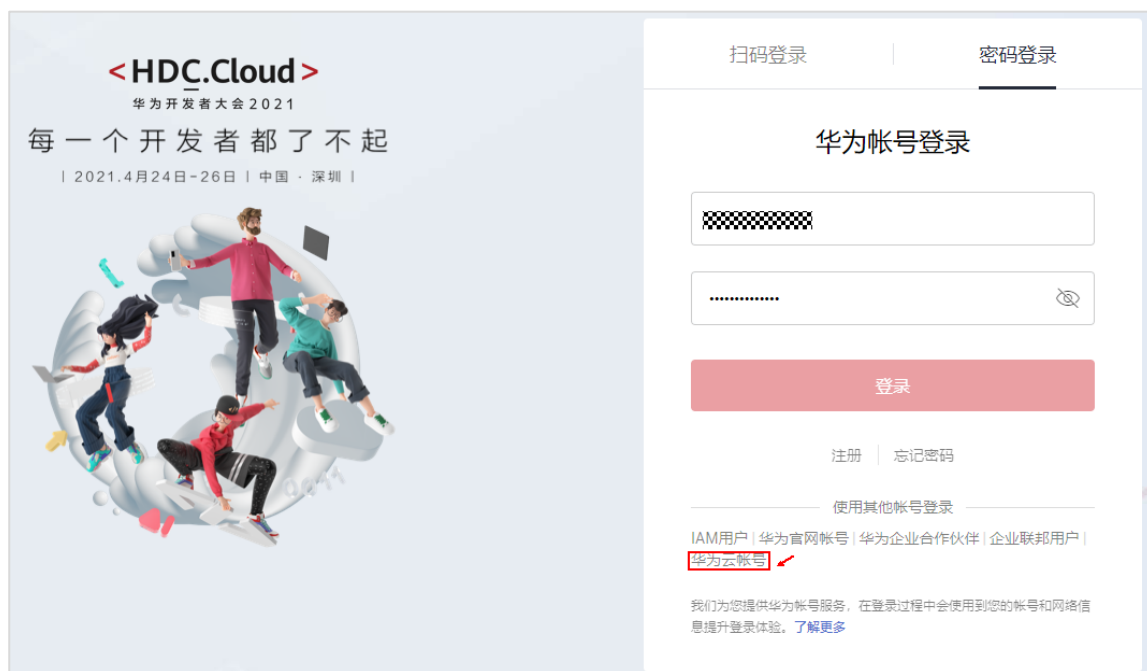
三：实验环境配置

步骤 1 登录华为云网站

<https://www.huaweicloud.com>



点击右上角登录，输入账号和密码



注意：华为云已统一登录入口，若仍不能登录则点击下方华为云账号进行登录。

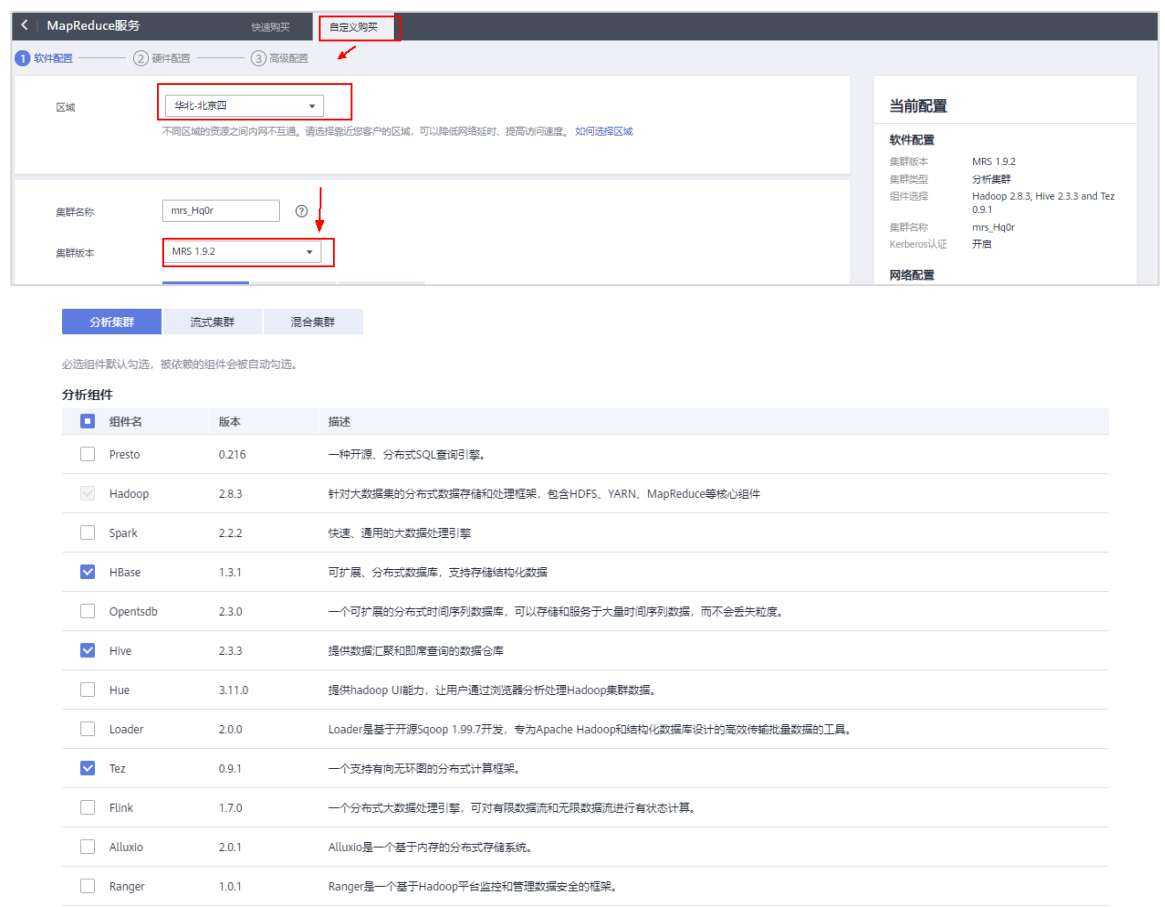
步骤 2 点击“EI 企业智能”选择“MapReduce 服务”



步骤 3 点击“立即购买”



选择“自定义购买” (这里还要选上 Spark)



点击下一步，进入硬件配置

选择“按需计费”，“可用区 2”，点击“弹性公网 IP”，如下图：



点击“购买弹性公网 IP”，选择“按需计费”，“按流量计费”，“5M”，点击

“立即购买”，如下图：



点击“提交”，如下图：

产品类型	产品规格	计费模式	数量	价格	
弹性公网IP	区域	北京四			
	类型	全动态BGP			
	IPv6转换	停用	按需计费	1	¥0.02/小时
带宽	带宽名称	bandwidth-818e			
	带宽类型	独享带宽			
	计费方式	按流量计费	按需计费	1	¥0.80/GB
	带宽大小	5 Mbit/s			

弹性公网IP费用 ¥0.02/小时 + 公网流量费用 ¥0.80/GB

上一步 提交

购买成功，如下图：

弹性公网IP

创建弹性公网IP使用体验调研，您宝贵的意见和建议是我们持续提升产品体验的动力，感谢您的参与！

解绑

修改带宽

续费

更多

所有状态

弹性公网IP

<input type="checkbox"/>	弹性公网IP	监控	状态	类型	带宽	带宽详情	已绑定实例	计费模式
<input type="checkbox"/>	39.9.141.144		未绑定	全动态BGP	--	--	--	按周 2021/04/21 15:22:16 创建

返回 MapReduce 服务自定义购买界面绑定 EIP

计费模式：包年/包月 按需计费

可用区：可用区1 可用区2 可用区3 可用区7

虚拟私有云：vpc-default

子网：subnet-dde5(192.168.0.0/24)

安全组：自动创建

弹性公网IP：暂不绑定

39.9.141.144

当前配置

软件配置

集群版本：MRS 1.9.2

集群类型：分析集群

组件选择：Hadoop 2.8.3 and HBase 1.2

集群名称：mrs_9xoh

Kerberos认证：开启

网络配置

计费模式：按需计费

区域：华北-北京四

可用区：可用区2

虚拟私有云：vpc-default

子网：subnet-dde5

安全组：自动创建

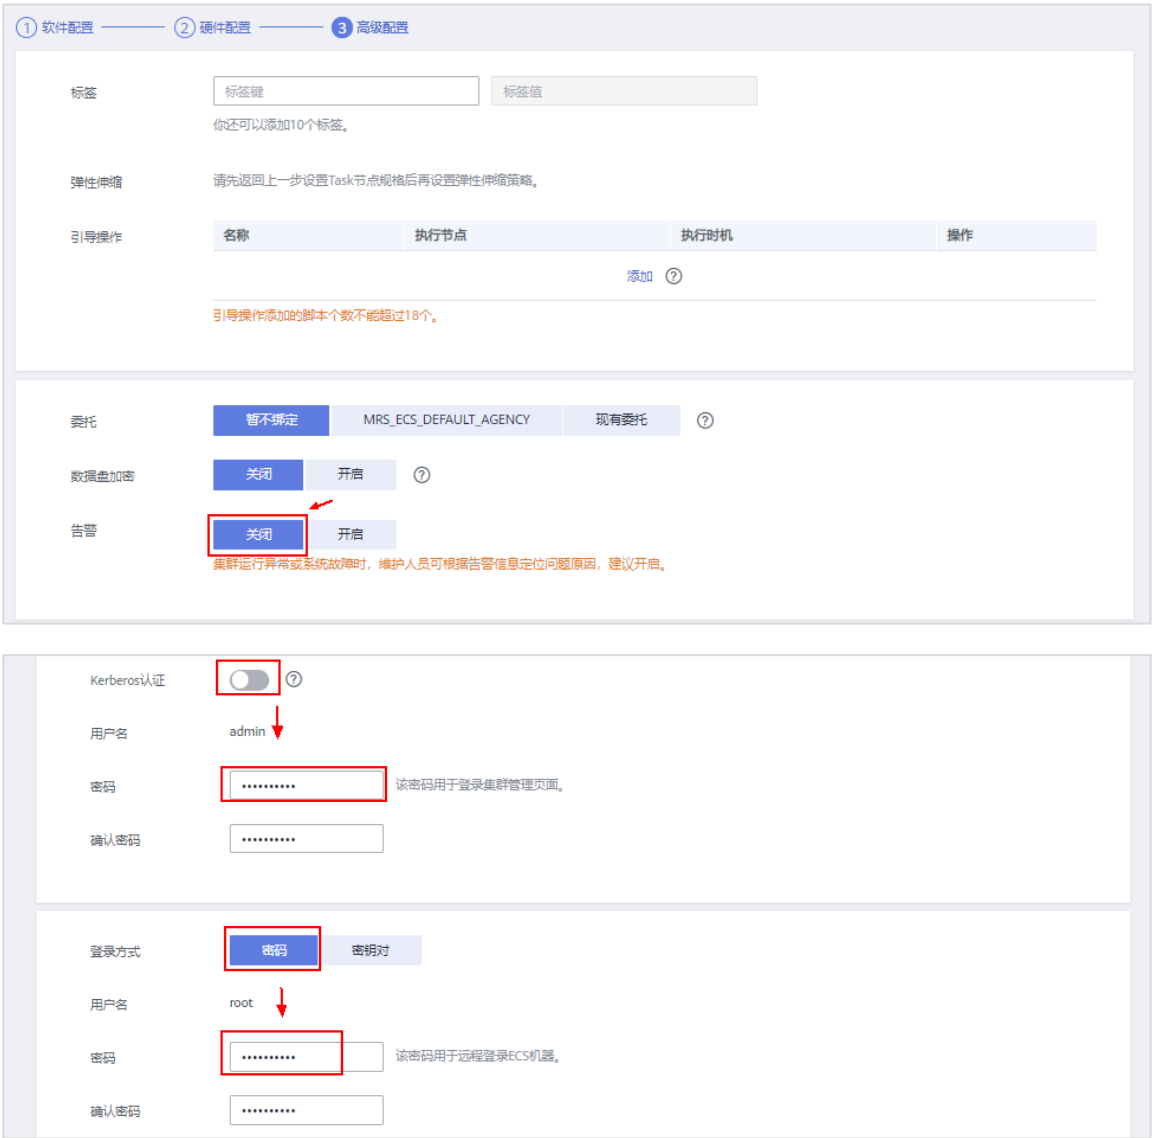
Master节点

选择“鲲鹏计算”，关闭高可用，调整 core 节点数为 1,如下图：



点击“下一步”

高级配置项参考如下：

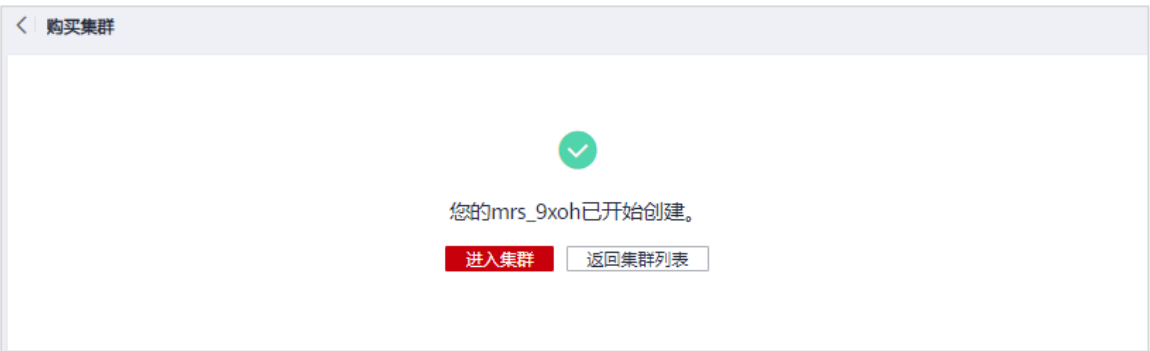


点击“确认授权”



点击“立即购买”

步骤 4 点击“返回集群列表”，如下图：



创建过程需要等待几分钟，待状态变为“运行中”集群创建完成

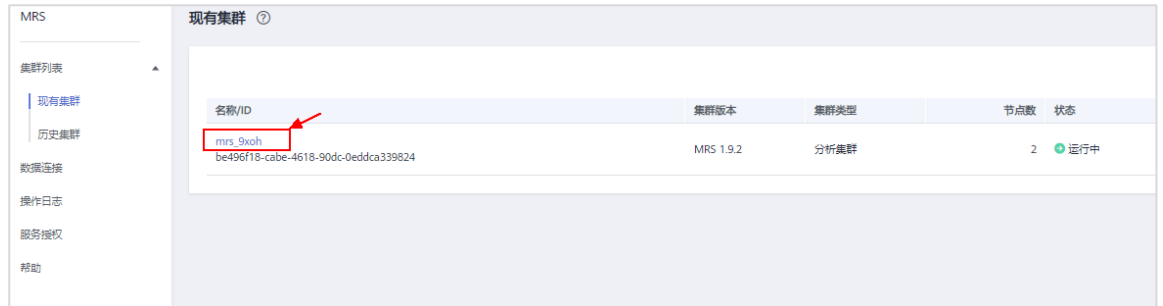


步骤 5 点击集群名称



步骤 6 配置安全组

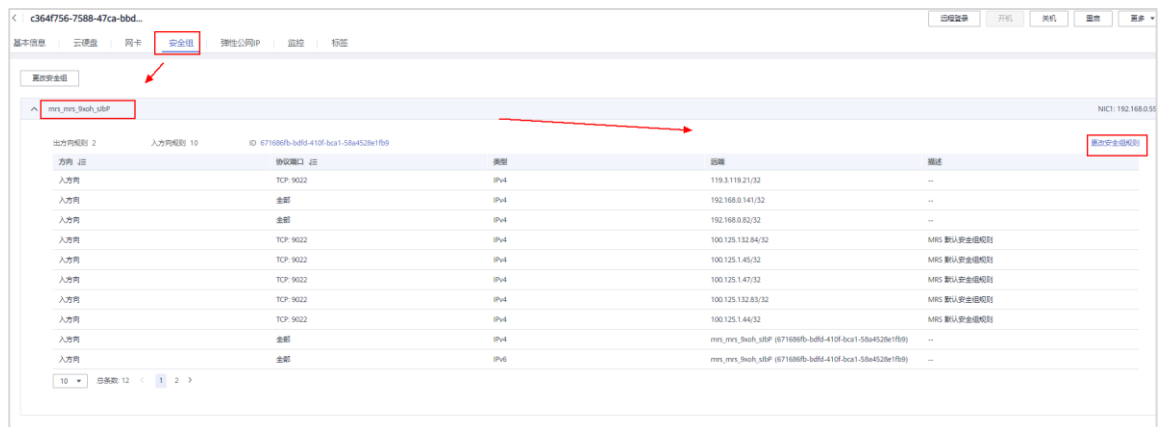
点击集群名称



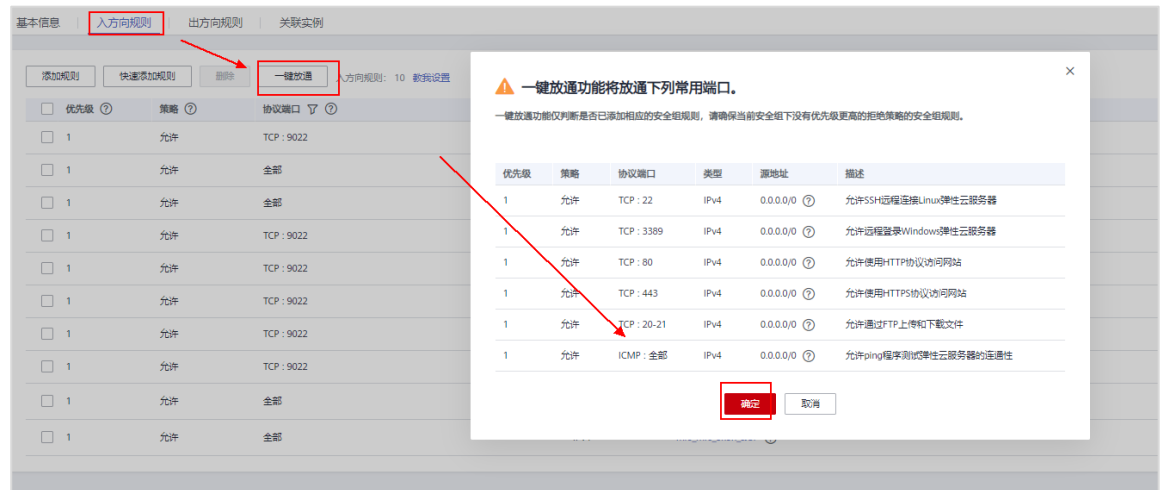
选择“节点管理”，点击含有“master1”的节点



在弹出页面中选择“安全组”，点击“更改安全组规则”，如下图所示：



选择“入方向规则”，点击“一键放通”，确认即可。

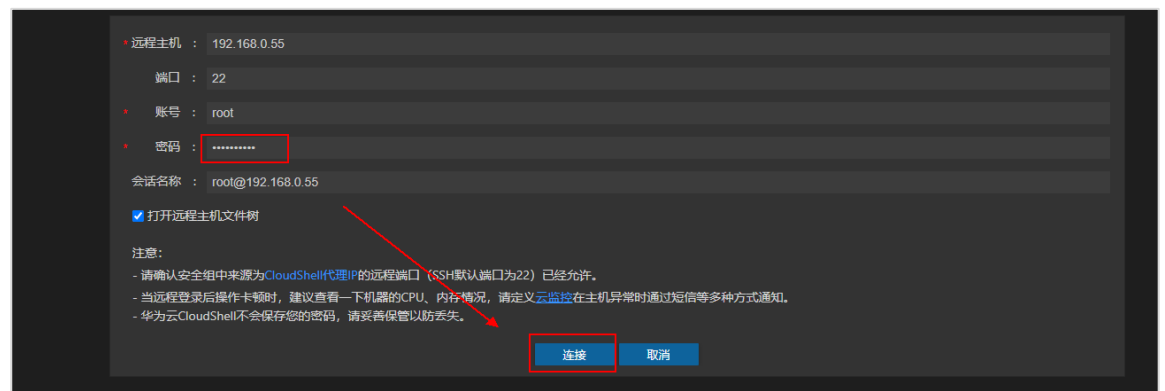


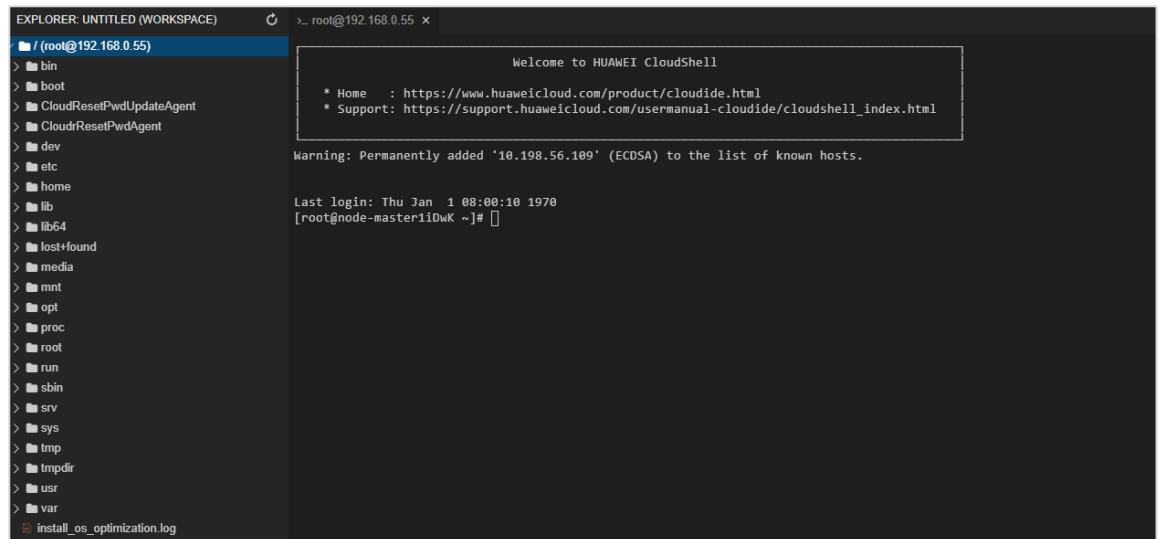
步骤 7 远程登录 master 节点

在安全组配置项，点击右上方“远程登录”，选择 cloudshell 登录。



输入密码，点击连接即可。





四：实验内容及步骤、实验的详细记录、实验结果分析

请附上实验过程截图（截图需包含指令）以及必要的文字分析

4.1 MapReduce

4.1.1 进入 hadoop(5')

`cd /opt/client/HDFS/Hadoop`

```
[root@node-master1jgZo ~]# cd /opt/client/HDFS/hadoop
[root@node-master1jgZo hadoop]#
```

4.1.2 添加环境变量(10')

`export HADOOP="/opt/client/HDFS/hadoop/share/hadoop"`

export

CLASSPATH="\$HADOOP/common/hadoop-common-2.8.3-mrs-1.9.0.jar:\$HADOOP/mapreduce/hadoop-mapreduce-client-core-2.8.3-mrs-1.9.0.jar:\$HADOOP/common/lib/commons-cli-1.2.jar:\$CLASSPATH"

```
[root@node-master1jgZo hadoop]# export HADOOP="/opt/client/HDFS/hadoop/share/hadoop"
[root@node-master1jgZo hadoop]# export CLASSPATH="$HADOOP/common/hadoop-common-2.8.3-mrs-1.9.0.jar:$HADOOP/mapreduce/hadoop-mapreduce-client-core-2.8.3-mrs-1.9.0.jar:$HADOOP/common/lib/commons-cli-1.2.jar:$CLASSPATH"
[root@node-master1jgZo hadoop]#
```

配置 java 环境

4.1.3 创建 java 程序 WordCount.java,在里面输入以下代码(5')

```
import java.io.IOException;
import java.util.StringTokenizer;

import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;

public class WordCount {

    public static class TokenizerMapper
        extends Mapper<Object, Text, Text, IntWritable>{

        private final static IntWritable one = new IntWritable(1);
        private Text word = new Text();

        public void map(Object key, Text value, Context context
            ) throws IOException, InterruptedException {
```

```

        StringTokenizer itr = new StringTokenizer(value.toString());
        while (itr.hasMoreTokens()) {
            word.set(itr.nextToken());
            context.write(word, one);
        }
    }
}

public static class IntSumReducer
    extends Reducer<Text,IntWritable,Text,IntWritable> {
    private IntWritable result = new IntWritable();

    public void reduce(Text key, Iterable<IntWritable> values,
                      Context context
                      ) throws IOException, InterruptedException {

        int sum = 0;
        for (IntWritable val : values) {
            sum += val.get();
        }
        result.set(sum);
        context.write(key, result);
    }
}

public static void main(String[] args) throws Exception {
    Configuration conf = new Configuration();
    Job job = Job.getInstance(conf, "word count");
    job.setJarByClass(WordCount.class);
    job.setMapperClass(TokenizerMapper.class);
    job.setCombinerClass(IntSumReducer.class);
    job.setReducerClass(IntSumReducer.class);
    job.setOutputKeyClass(Text.class);
    job.setOutputValueClass(IntWritable.class);
    FileInputFormat.addInputPath(job, new Path(args[0]));
    FileOutputFormat.setOutputPath(job, new Path(args[1]));
    System.exit(job.waitForCompletion(true) ? 0 : 1);
}
}

```

```
[root@node-master1dlKK hadoop]# vim WordCount.java
[root@node-master1dlKK hadoop]#
```

```
import java.io.IOException;
import java.util.StringTokenizer;

import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;

public class WordCount {
    public static class TokenizerMapper
        extends Mapper<Object, Text, Text, IntWritable>{
        private final static IntWritable one = new IntWritable(1);
        private Text word = new Text();

        public void map(Object key, Text value, Context context)
            throws IOException, InterruptedException {
            StringTokenizer itr = new StringTokenizer(value.toString());
            while (itr.hasMoreTokens()) {
                word.set(itr.nextToken());
                context.write(word, one);
            }
        }
    }

    public static class IntSumReducer
        extends Reducer<Text, IntWritable, Text, IntWritable> {
        private IntWritable result = new IntWritable();

        public void reduce(Text key, Iterable<IntWritable> values,
            Context context)
            throws IOException, InterruptedException {
            for (IntWritable val : values) {
                result.add(val.get());
            }
            context.write(key, result);
        }
    }
}
```

编写 java 文件

4.1.4 编译 WordCount.java(5')

```
[root@node-master1dlKK hadoop]# javac WordCount.java
[root@node-master1dlKK hadoop]# ls
bin          LICENSE.txt  'WordCount$IntSumReducer.class'
etc          NOTICE.txt  'WordCount$TokenizerMapper.class'
hdfs-c-example README.txt   WordCount.class
include      sbin        WordCount.java
lib          share
libexec     version.properties
[root@node-master1dlKK hadoop]#
```

```
[root@node-master1dlKK hadoop]# jar cf WordCount.jar WordCount*.class
[root@node-master1dlKK hadoop]# ls
bin          NOTICE.txt  'WordCount$IntSumReducer.class'
etc          README.txt  'WordCount$TokenizerMapper.class'
hdfs-c-example sbin        WordCount.class
include      share       WordCount.jar
lib          test1       WordCount.java
libexec     test2
LICENSE.txt  version.properties
[root@node-master1dlKK hadoop]#
```

4.1.5 创建文件 test1, 内容为 hello hust, 文件 test2, 内容为 hello 学号, 将他们放入 hdfs 的 /input 文件夹内。(方法见实验一)(5')

```
[root@node-master1dlKK hadoop]# vim test1
[root@node-master1dlKK hadoop]# vim test2
[root@node-master1dlKK hadoop]#
```

```
hello hust
```

```
~
~
~
~
```

```
hello X2020I1007
```

```
~
~
~
~
```

```
[root@node-master1dlKK hadoop]# hdfs dfs -mkdir /input
[root@node-master1dlKK hadoop]# hdfs dfs -put test1 /input
[root@node-master1dlKK hadoop]# hdfs dfs -put test2 /input
[root@node-master1dlKK hadoop]# hdfs dfs -ls /input
Found 2 items
-rw-r--r--    1 root ficommon      11 2021-12-14 20:27 /input/test1
-rw-r--r--    1 root ficommon      17 2021-12-14 20:28 /input/test2
[root@node-master1dlKK hadoop]#
```

4.1.6 运行 WordCount.jar 将 hdfs 的 /input 作为输入，/output 作为输出，并打印 /output 目录下的文件，显示出词频统计的结果(5')

export

HADOOP_CLASSPATH=\$HADOOP_CLASSPATH:/opt/client/HDFS/hadoop/WordCount.jar

hadoop jar WordCount.jar WordCount hdfs:///input hdfs:///output

hdfs dfs -cat /output/part-r-00000

```
[root@node-master1dlKK hadoop]# export HADOOP_CLASSPATH=$HADOOP_CLASSPATH:/opt/client/HDFS/hadoop/WordCount.jar
[root@node-master1dlKK hadoop]# hadoop jar WordCount.jar WordCount hdfs:///input hdfs:///output
WARNING: Use "yarn jar" to launch YARN applications.
21/12/14 20:32:39 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.
21/12/14 20:32:40 INFO input.FileInputFormat: Total input files to process : 2
21/12/14 20:32:42 INFO mapreduce.JobSubmitter: number of splits:2
21/12/14 20:32:43 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1639483812097_0002
21/12/14 20:32:44 INFO impl.YarnClientImpl: Submitted application application_1639483812097_0002
21/12/14 20:32:44 INFO mapreduce.Job: The url to track the job: http://node-master1dlkk.mrs-fbva.com:8088/proxy/application_1639483812097_0002/
21/12/14 20:32:44 INFO mapreduce.Job: Running job: job_1639483812097_0002
[root@node-master1dlKK hadoop]# hdfs dfs -cat /output/part-r-00000
X2020I1007    1
hello        2
hust         1
[root@node-master1dlKK hadoop]#
```



```
>>> counts = lines.flatMap(lambda x: x.split(' ')).map(lambda x: (x, 1)).reduceByKey(lambda x, y: x + y)
>>> output = counts.collect()
2021-12-14 20:41:04,911 | INFO | Thread-3 | Starting job: collect at <stdin>:1 | org.apache.spark.internal.Logging$class.logInfo(Logging.scala:54)
2021-12-14 20:41:04,934 | INFO | dag-scheduler-event-loop | Registering RDD 5 (reduceByKey at <stdin>:1) | org.apache.spark.internal.Logging$class.logInfo(Logging.scala:54)
2021-12-14 20:41:04,937 | INFO | dag-scheduler-event-loop | Got job 0 (collect at <stdin>:1) with 2 output partitions | org.apache.spark.internal.Logging$class.logInfo(Logging.scala:54)
2021-12-14 20:41:04,938 | INFO | dag-scheduler-event-loop | Final stage: ResultStage 1 (collect at <stdin>:1) | org.apache.spark.internal.Logging$class.logInfo(Logging.scala:54)
2021-12-14 20:41:04,938 | INFO | dag-scheduler-event-loop | Parents of final stage: List(ShuffleMapStage 0) | org.apache.spark.internal.Logging$class.logInfo(Logging.scala:54)
2021-12-14 20:41:04,941 | INFO | dag-scheduler-event-loop | Missing parents: List(ShuffleMapStage 0) | org.apache.spark.internal.Logging$class.logInfo(Logging.scala:54)
2021-12-14 20:41:04,947 | INFO | dag-scheduler-event-loop | Submitting ShuffleMapStage 0 (PairwiseRDD[5] at reduceByKey at <stdin>:1), which has no missing parents | org.apache.spark.internal.Logging$class.logInfo(Logging.scala:54)
2021-12-14 20:41:05,054 | INFO | dag-scheduler-event-loop | Block broadcast_1 stored as values in memory (estimated size 13.0 KB, free 92.8 MB) | org.apache.spark.internal.Logging$class.logInfo(Logging.scala:54)
2021-12-14 20:41:05,080 | INFO | dag-scheduler-event-loop | Block broadcast_1 piece0 stored as bytes in memory (estimated size 8.0 KB, free 92.8 MB) | org.apache.spark.internal.Logging$class.logInfo(Logging.scala:54)
2021-12-14 20:41:05,087 | INFO | dispatcher-event-loop-3 | Added broadcast_1 piece0 in memory on *.*.0.142:22873 (size: 8.0 KB, free: 93.3 MB) | org.apache.spark.internal.Logging$class.logInfo(Logging.scala:54)
```

4.2.4 输出结果(10')

```
>>> print(output)
[(u'hust', 1), (u'X2020I1007', 1), (u'hello', 2)]
>>> █
```

4.3 附加题

file1:

20210001 Math 90

20210002 Math 80

20210003 Math 70

file2:

20210001 English 80

20210002 English 70

20210003 English 60

1.将以上两个文件存入 hdfs(10')

```
[root@node-master1dlKK ~]# vim file1
[root@node-master1dlKK ~]# vim file2
[root@node-master1dlKK ~]#
```

20210001 Math 90	20210001 English 80
20210002 Math 80	20210002 English 70
20210003 Math 70	20210003 English 60
~	~
~	~
~	~
~	~
~	~

```
[root@node-master1dlKK ~]# hdfs dfs -mkdir /hdfs
[root@node-master1dlKK ~]# hdfs dfs -put file1 /hdfs
[root@node-master1dlKK ~]# hdfs dfs -put file2 /hdfs
[root@node-master1dlKK ~]# hdfs dfs -ls /hdfs
Found 2 items
-rw-r--r--  1 root ficommon      52 2021-12-14 20:56 /hdfs/file1
-rw-r--r--  1 root ficommon      61 2021-12-14 20:56 /hdfs/file2
[root@node-master1dlKK ~]#
```

将上述内容编写入文件并存入名为 hdfs 文件夹内

2.编写 mapreduce 的程序，输出每门课的平均成绩。(10')

3.编写 mapreduce 的程序，输出每位同学有多少门课成绩低于 75 分。(10')

五：实验总结(10')

通过本次实验，让我学会了在 Hadoop 平台上配置 java 环境并且编译.java 文件对数据文件 test1 与 test2 进行处理以及在 Spark 平台上使用 python 代码对数据文件进行词频统计计算，从而实现 WordCount 功能（单词词频统计计算功能）。从中对 MapReduce 大数据并行处理的实现方法有了更多的认知。而本次最大的难点在于我还没学过 Java 这门语言，以至于在附加题最后两题上出现了卡顿，不知道如何实现那两个算法，同时还读不懂 Java 程序。因此我认识到了自己的短板在于缺少了 Java 编程语言方面的知识，在后面的实验开始前，我会好好地恶补一下这方面的知识，从而使实验完成得更加顺利。

附录：

1 MapReduce 官方教程

<https://hadoop.apache.org/docs/r2.8.3/hadoop-mapreduce-client/hadoop-mapreduce-client-core/MapReduceTutorial.html>

2 Spark 官方教程

<https://spark.apache.org/examples.html>